

Shallow Water: Initial Report

Group Number: 2

Members: Nimit Sohoni (nss66), Edward Tremel (ejt64), Ian Vermeulen (iyv2)

Optimizations used or attempted:

Restructuring and annotation to enable greater vectorization, informed by results from profiling
Parallelization with OpenMP

More details:

1. Vectorization
 - a. Used compiler directives such as '#pragma ivdep' to encourage the compiler to vectorize loops and ignore possible dependencies that did not actually exist.
 - b. Attempted to restructure the code to make better use of memory locality and vectorization.
2. Parallelization
 - a. Split the simulation grid into rectangular blocks based on the number of threads available.
 - i. For an a-by-b grid with p threads, the block size would be $\text{floor}(a/\sqrt{p})$ -by- $\text{floor}(b/\sqrt{p})$. Some cells may not be covered by any blocks, so those are processed separately. However, the number of these cells should be small relative to the entire grid size, so should not affect performance too much.
 - ii. We will also have k layers of ghost cells for each block, so that each thread can compute k time steps before communication is necessary. k is a parameter that we will vary to figure out which value gives the best performance.

Results:

When we ran the `amplxe-cl` profiling tool on the default `dam_break` test case, we found that by far the bottleneck was the `limited_derivs` function.

amplxe: Executing actions 50 % Generating a report	Function	Module	CPU Time	Spin Time	Overhead Time
-----	-----	-----	-----	-----	-----
Central2D<Shallow2D, MinMod<float>>::limited_derivs	shallow		1.148s	0s	0s
Central2D<Shallow2D, MinMod<float>>::compute_step	shallow		0.659s	0s	0s
Central2D<Shallow2D, MinMod<float>>::compute_fg_speeds	shallow		0.230s	0s	0s
[Outside any known module]	[Unknown]		0.021s	0s	0s
_IO_file_xsputn	libc-2.12.so		0.015s	0s	0s
_IO_fwrite	libc-2.12.so		0.010s	0s	0s
Central2D<Shallow2D, MinMod<float>>::solution_check	shallow		0.004s	0s	0s
SimViz<Central2D<Shallow2D, MinMod<float>>>::write_frame	shallow		0.004s	0s	0s
Central2D<Shallow2D, MinMod<float>>::Central2D	shallow		0.003s	0s	0s
Central2D<Shallow2D, MinMod<float>>::run	shallow		0.001s	0s	0s
std::array<float, (unsigned long)3>::operator[]	shallow		0.001s	0s	0s
Central2D<Shallow2D, MinMod<float>>::offset	shallow		0.001s	0s	0s
do_lookup_x	ld-2.12.so		0.001s	0s	0s

After we made some changes to improve vectorization, we were able to speed up the `compute_step` and `compute_fg_speeds` functions slightly:

amplxe: Executing actions 50 % Generating a report	Function	Module	CPU Time	Spin Time	Overhead Time
Central2D<Shallow2D, MinMod<float>>>::limited_derivs	shallow		1.149s	0s	0s
Central2D<Shallow2D, MinMod<float>>>::compute_step	shallow		0.548s	0s	0s
Central2D<Shallow2D, MinMod<float>>>::compute_fg_speeds	shallow		0.130s	0s	0s
[Outside any known module]	[Unknown]		0.017s	0s	0s
_IO_file_xsputn	libc-2.12.so		0.012s	0s	0s
_IO_fwrite	libc-2.12.so		0.010s	0s	0s
Central2D<Shallow2D, MinMod<float>>>::solution_check	shallow		0.005s	0s	0s
SimViz<Central2D<Shallow2D, MinMod<float>>>::write_frame	shallow		0.005s	0s	0s
Central2D<Shallow2D, MinMod<float>>>::Central2D	shallow		0.002s	0s	0s
Central2D<Shallow2D, MinMod<float>>>::run	shallow		0.002s	0s	0s
std::array<float, (unsigned long)3>::operator[]	shallow		0.002s	0s	0s
Central2D<Shallow2D, MinMod<float>>>::offset	shallow		0.001s	0s	0s
do_lookup_x	ld-2.12.so		0.001s	0s	0s

As can be seen, both the `compute_step` and `compute_fg_speeds` functions take about 0.1 seconds less time after vectorization is used. These numbers were not just coincidences or artifacts, but held up after repeated trials.

The `#pragma ivdep` directive was helpful in loops where vectorization was only being held back by the compiler's prior inability to assume that the memory locations being modified and the memory locations being read did not overlap. However, in many cases, such as the `limited_derivs` function, the loops still did not vectorize, because they would actually have become slower than the serial version due to memory access patterns that disrupted cache locality. Therefore, it would be necessary to change the structure of the code to enable access patterns making use of cache locality. For instance, instead of using vectors of 'vecs' for `u`, `f`, `g`, and so on, we could instead store all this data in a flattened (single-dimensional) data structure and change our 'offset' function to convert three input numbers into an offset within this array, similar to what Professor Bindel did in his C code. We may use `valarrays` for this purpose. Indeed, for the purposes of this submission, we opted to instead go with the C code provided by Professor Bindel, to benchmark our parallelization, and meanwhile simultaneously another team member would work on getting the C++ code properly vectorized.

Parallelization results:

Unfortunately, we were unable to finish implementing the parallelization before the assignment was due.

Future steps:

We will finish restructuring the C++ code to attempt to get it close to the performance of the C code. We will finish implementing the grid parallelization scheme, and will tune parameters to find out the optimal number of layers of ghost cells per block.