

CS 5220

Project 2 - Shallow Water Simulation

Marc Aurele Gilles (mtg79)
Sheroze Sherifdeen(mss385)

October 18, 2015

1 Introduction

Define structured grid computations
Define shallow water simulations
Math overview

2 Design Decisions

The following sections describe the implementation changes from the original code found at <https://github.com/cornell-cs5220-f15/water>.

2.1 Memory Layout

The original solution a two dimensional vectors of 3-vectors to represent U_t , $F(U)_x$, and $G(U)_y$. During each time step, the solution accesses each element in the 2-D grid sequentially. Then, the `vector<vector<real>>` representation leads to memory accesses that are not local spatially.

Therefore, our solution chooses to use 3 separate two dimensional vectors per objects, U_t , $F(U)_x$, and $G(U)_y$. Thus, we are required in general to perform three loop iterations in place of a single loop in the original solution. But this approach leverages spatial locality, especially in `compute_step` and `limited_derivs` functions.

2.2 Vectorization

By observing the profiling information, we noticed that the original solution spends majority of its computational time in the functions `limited_derivs`, `compute_step` and `compute_fg_speeds`. By adopting the newer memory layout, we enabled spatially local memory accesses. We were also able to decompose `for` loops in the solution to improve vectorization. Refer to `ipo_out_vectorization.optrpt` in <https://github.com/sheroze1123/water/tree/vectorization> for more information.

In `compute_fg_speeds`, we performed two separate loops to compute flux and wave speeds. The flux computation and the wave speed computation for the complete grid is not handled by the `Physics` class. We used `#pragma simd` directives to instruct the compiler the ability to vectorize

these computations. The compiler was successfully able to vectorize these functions with an estimated potential speedup of 6.7.

`limited_derivs` uses the `limdiff` function in `minmod.h`. To improve the vectorization of this computation, we changed the implementation of `limdiff` in the following ways.

1. `limdiff` now performs the computation on the complete grid instead of at one grid point.
2. `limdiff` was decomposed as `limdiff_x` and `limdiff_y` to perform the limiter along the x dimension and the y dimension separately while still retaining unit stride.

2.3 Parallelization

2.4 Domain Decomposition

3 Analysis

3.1 Profiling

3.1.1 Original solution

We began the optimization by analyzing the time profiles of the original code.

Function	Module	CPU Time	Spin Time	Overhead Time
Central2D<Shallow2D, MinMod<float>>::limited_derivs	shallow	2.529 s	0 s	0 s
Central2D<Shallow2D, MinMod<float>>::compute_step	shallow	1.210 s	0 s	0 s
Central2D<Shallow2D, MinMod<float>>::compute_fg_speeds	shallow	0.426 s	0 s	0 s
_IO_file_xsputn	libc-2.12.so	0.027 s	0 s	0 s
_IO_fwrite	libc-2.12.so	0.025 s	0 s	0 s

3.2 Speedup Plots

References

- [1] Data Alignment to Assist Vectorization. (n.d.). Retrieved September 30, 2015, from <https://software.intel.com/en-us/articles/data-alignment-to-assist-vectorization>