

Homework 2

CS 5220

Lara Backer, Xiang Long, Saul Toscano

October 19, 2015

Initial Profiling - Vtune

An initial profiling of the shallow wave code on the Totient cluster was done using Vtune, to identify which regions of the code took the longest time to run. This was done in order to target parallelization of the functions that took the longest time to complete.

Figure 1 displays the Vtune assessment of the overall program. The top three most time consuming functions are limited_derivs, compute_step, and compute_fg_speeds.

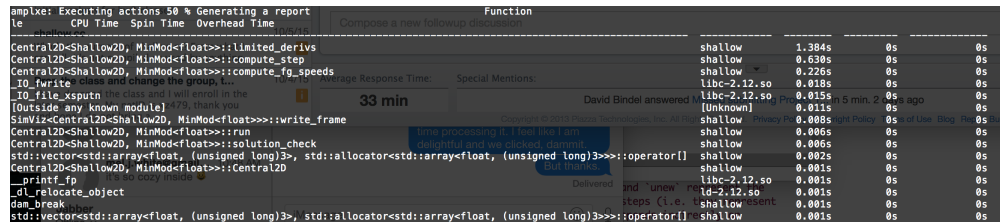


Figure 1: Most time consuming functions for shallow wave code on Totient cluster

The function limited_derivs is used to limit the estimation of the derivative of the fluxes by calling the limdiff function. Compute_fg_speeds is used to evaluate the fluxes at the cell centers. The compute_step function is the seciton of the code that evaluates both the predictor and corrector calculations for the timestep. To see the time spent in each step using Vtune, the specific functions must be referenced. While this is not overly useful for the limited_derivs function, which only performs the two limdiff calls, the in-depth timings for the other two most expensive functions are shown in Figures 2 and 3.

Vectorization

One method of speeding up the code is to apply vectorization, applying operations to arrays instead of array sub elements, reducing the number of required computations. To apply vectorization to our code for instance, we grouped the fluxes into single arrays, specifying pointers to the locations of each flux, and indicating the strides required for each component.

The .optreport file output shows the sections that could be vectorized to improve efficiency.

openMP Parallelization

OpenMP can be used to further improve performance by multithreading, where each thread executes a section of the code independently. In C++, OpenMP uses #pragmas to signify the parallelization of code. The pragma omp parallel, for instance, is used to fork threads to carry out the enclosed work. Additional pragmas can be used to specify sections of work, or splitting up loop iterations.

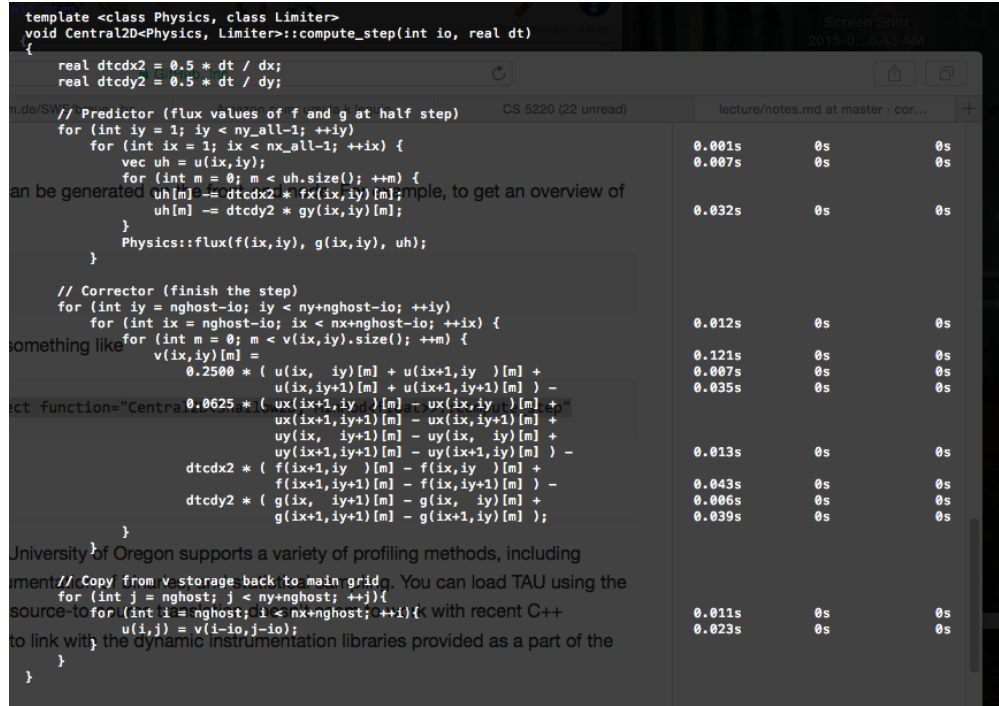


Figure 2: compute step function timings

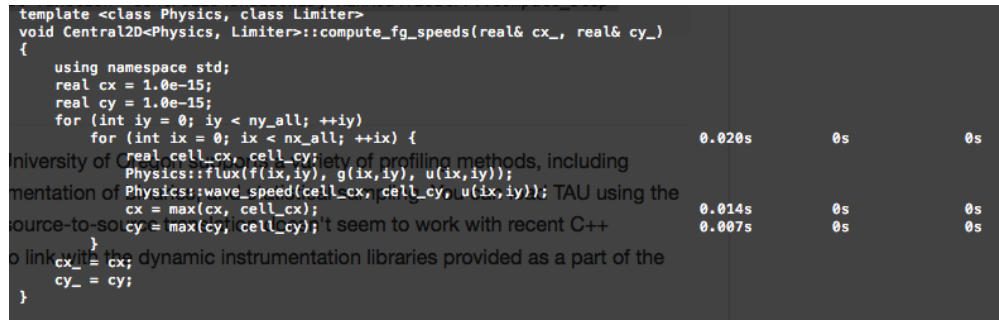


Figure 3: fgspeds function timings

Domain Decomposition

The method best suited to parallelize the shallow water problem is domain decomposition. For this method, the overall domain is split and solved for on separate processors. Communications between processors are needed to solve for ‘ghost cells’ that overlap neighboring domains. However, communications can overtake the domain computations in cost, depending on the number of processors and communications compared to the domain size that each processor has to solve, so scaling trials must be employed to determine the optimal problem size per processor.